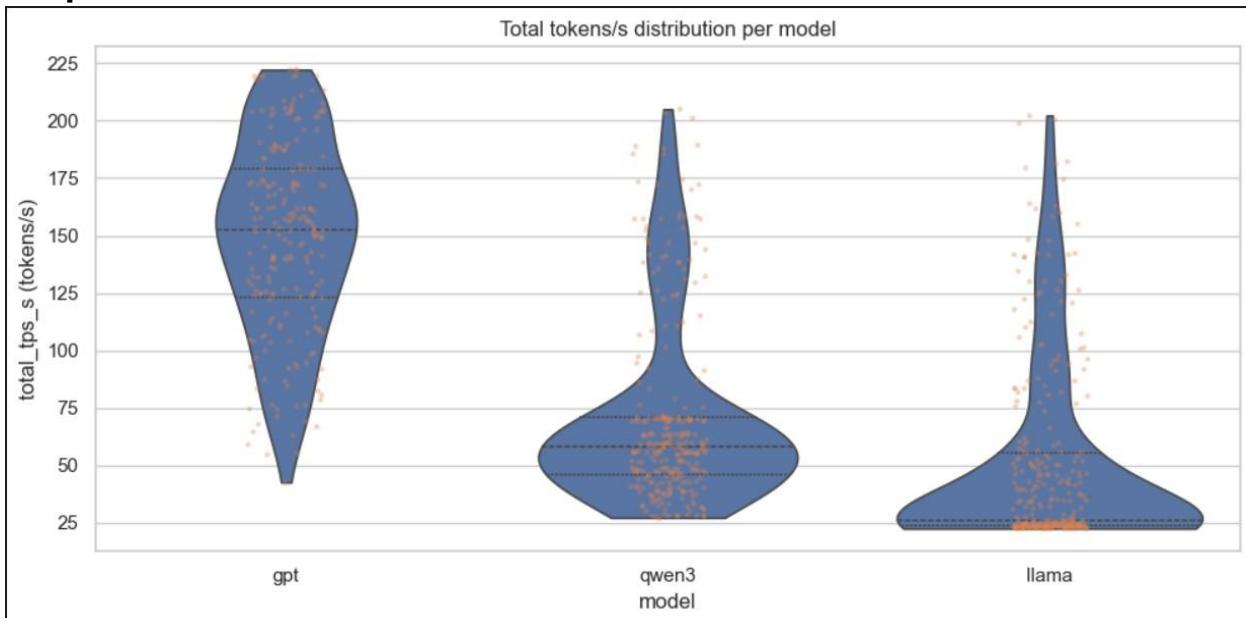


# Graph Analysis and Optimization Directions

## Graph A : Token/s Distribution



### Graph analysis

- GPT: very high throughput.
- Qwen3: medium typical throughput, with a few spikes.
- Llama: low typical throughput, with a long tail (rare spikes).

### Likely vulnerabilities

#### 1 Instability / high variance (especially Qwen3 & Llama)

When a model is often low and sometimes very high, it usually indicates that performance depends heavily on external conditions (batching, queueing, sequence lengths, contention).

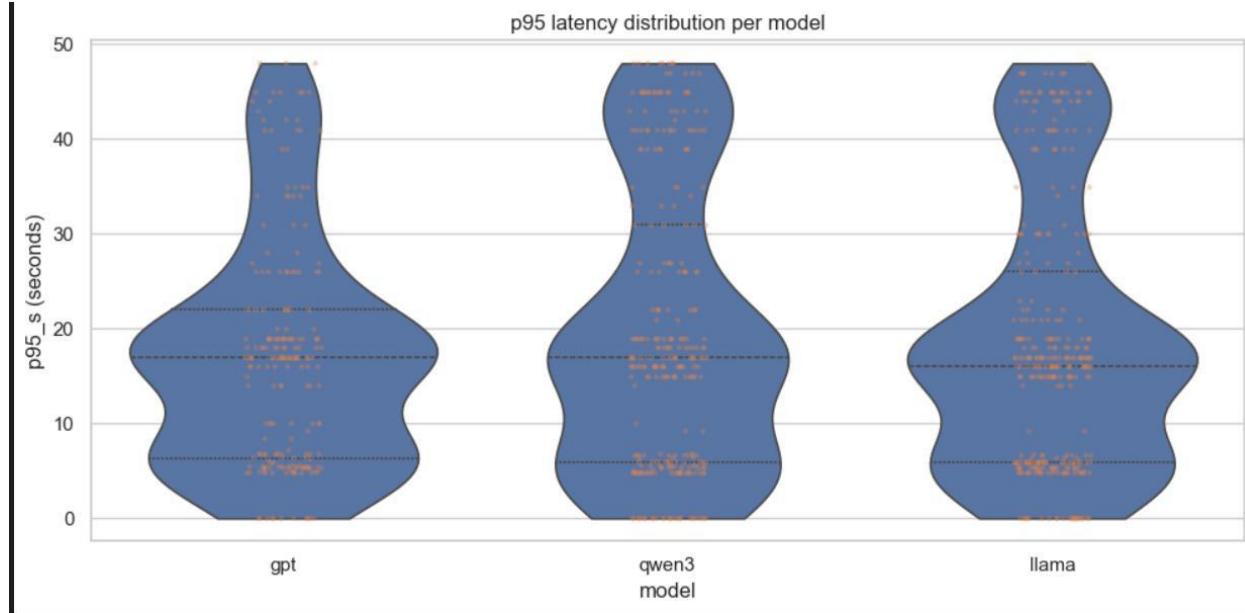
#### 2 Sensitivity to load conditions

“Long tails” mean the model can fall into a degraded regime (saturated GPU, scheduler waiting, less favorable batching).

### Plausible explanations

- **Variable output length:** observed tokens/s changes (longer decode means more stable; very short means noisy).
- **Batch size changes:** some moments benefit from a perfect batch, others do not.
- **Warm-up:** slower start, then stable regime.

## Graph B : p95 Distribution



### Graph analysis

- Typical p95 15-20s on all 3 models.
- Presence of a low mode (5-7s) and a high mode (40-45s).

### Likely vulnerabilities

#### 1 Queueing / congestion (p95 goes very high)

A p95 at 40-45s is typical of a queue or a bottleneck (requests are waiting).

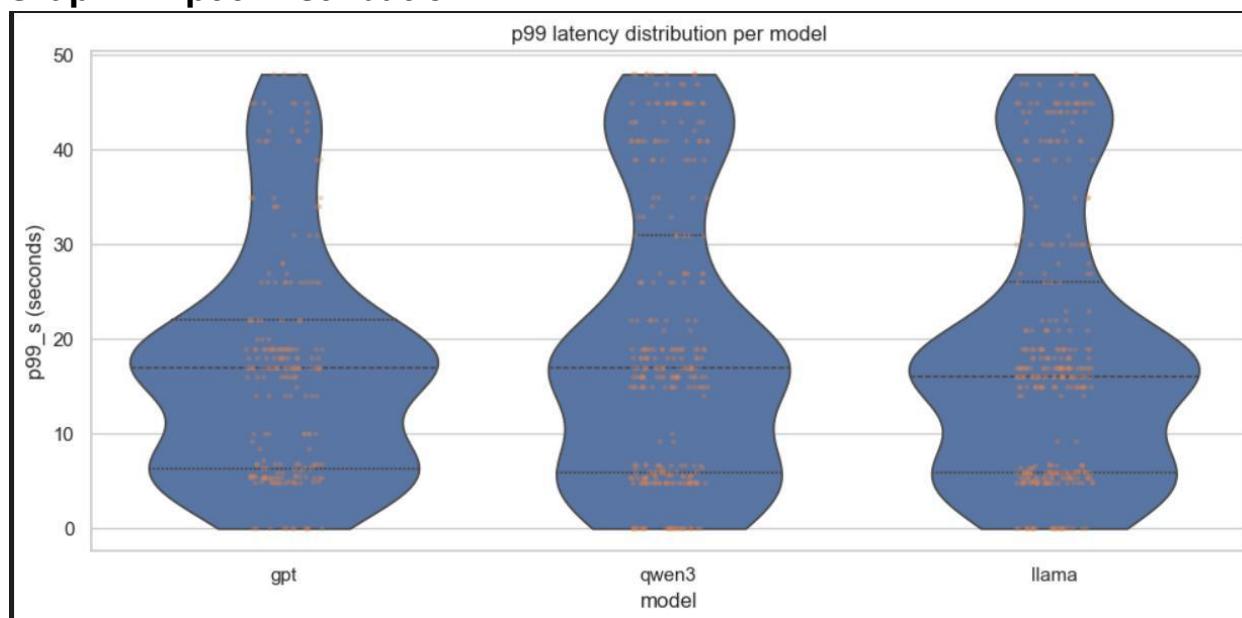
#### 2 Multiple regimes

The system switches between “normal” and “stress” (modes visible on the violin plot).

### Plausible explanations

- **GPU saturation** means `num_requests_waiting` would increase.
- **Prompt length / output length variability**: some requests are much more expensive.
- **Batching**: under load, a request may wait for the next batch.

## Graph C : p99 Distribution



### Graph analysis

- p99 is very close to p95, with the same spikes up to ~45–48s.

### Likely vulnerabilities

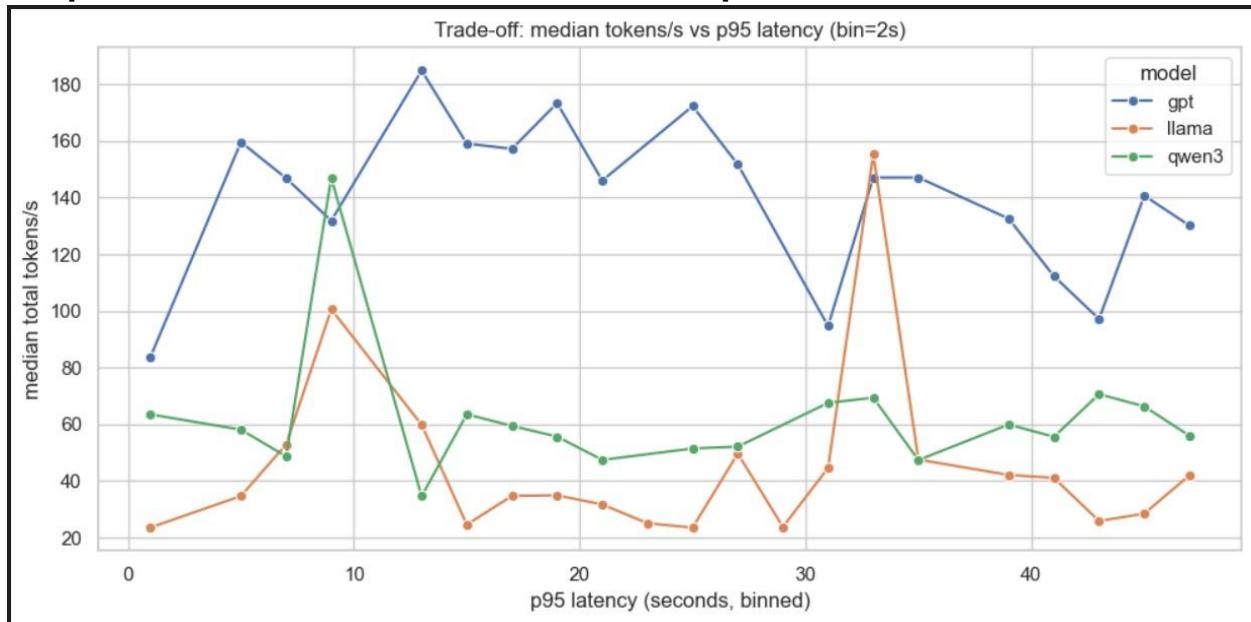
#### High and persistent worst-case

High p99 means a few requests regularly experience extreme delays.

### Plausible explanations

- Head-of-line blocking: large requests slow down others (batch / scheduler).

## Graph D : Trade-off: median tokens/s vs p95



## Graph analysis

- GPT keeps high throughput even when p95 increases.
- Qwen3 is stable but low.
- Llama is very low with one isolated peak.

## Likely vulnerabilities

### 1 Llama: frequent “degraded” regime

Low tokens/s across almost all bins, inefficiency risk (heavy model, config/quantization).

### 2 Qwen3: performance ceiling

It does not exploit favorable regimes as much as GPT (may be limited by config: KV cache, tensor parallel, quantization).

## Plausible explanations

- GPT might be served on a more optimized setup / more GPUs / better parallelization.
- Qwen3/Llama may be more sensitive to length (heavy prefill).

## **Parameters to tune per model (vLLM optimization hypotheses)**

**Important:** these settings are hypothetical. The goal is to target the observed symptoms: queueing, variability, throughput ceiling, and tail latency.

### **1) GPT (high throughput but p95/p99 latency sometimes high)**

**Goal:** reduce tail latency without killing throughput

- Reduce the queue (if load varies):
  - max-num-seqs (lower if too much concurrency to have fewer requests waiting).
  - max-num-batched-tokens (tune: too high can increase latency for short requests).
- Stabilize prefill:
  - enable-chunked-prefill (better responsiveness with long prompts).
- Limit context length if possible:
  - max-model-len (if you don't need a huge context).
- Tune GPU memory:
  - gpu-memory-utilization (too aggressive can create memory).

### **2) Qwen3 (medium throughput and spikes)**

**Goal:** increase typical throughput and reduce variance

- Improve batching efficiency (if many small requests):
  - max-num-batched-tokens (increase gradually).
  - max-num-seqs (increase if GPU under-utilized, decrease if queueing).
- Optimize multi-GPU parallelization (if the model is very large):
  - tensor-parallel-size (TP): test values that better saturate GPUs.
  - pipeline-parallel-size (PP): if the model is huge and TP alone is not enough.
- KV cache / memory :
  - kv-cache-dtype (FP8/FP16 depending on support) to reduce memory and increase capacity.
  - gpu-memory-utilization.
- Reduce prefill overhead:
  - enable-chunked-prefill (if prompts are long).
  - enable-prefix-caching

### 3) Llama (low typical throughput + rare spikes → frequent “degraded” regime)

**Goal:** get out of the low regime and improve the “typical” level

- Check/tune quantization:

--quantization (AWQ/GPTQ depending on the build): can significantly increase throughput, sometimes at a small quality cost.

- Parallelism:

--tensor-parallel-size / --pipeline-parallel-size

- Batching:

--max-num-batched-tokens (increase if Llama never benefits from good batches).

--max-num-seqs (tune to avoid exploding p95/p99).

- Reduce latency variability:

--enable-chunked-prefill

--max-model-len (if unnecessarily large, it costs memory and can hurt performance).

- Memory/KV :

--gpu-memory-utilization and --kv-cache-dtype to avoid situations where the cache is too constrained